

Consistency in Transactional Distributed Databases: Protocols and Testing

Hengfeng Wei (魏恒峰)


hfwei@nju.edu.cn

Nov. 04, 2022




Background

Centralized Databases *vs.* Distributed Databases



Data Consistency Problem

“Data Partition + Data Replication”




Data Consistency Problem

Data Consistency Problem

(Strong) Consistency, Availability, Latency, Partition tolerance





CAP Theorem



PACELC Tradeoff

Consistency Models

Use **consistency models** to capture these tradeoffs



Researches on data consistency around consistency models:

Computability: What is possible or impossible?

Protocol: How to design fast, scalable, and fault-tolerant protocols?

Testing: What is the complexity?
How to design efficient testing algorithms?

Researches on data consistency around consistency models:

Computability: What is possible or impossible?

Protocol: How to design fast, scalable, and fault-tolerant protocols?

Testing: What is the complexity?
How to design efficient testing algorithms?

Classic problems with the ever-changing requirements

My Researches

Research I (≥ 2012): Read/Write Register^a



Distributed NoSQL Key-Value Stores (TODO: 重新画图)
TODO: +research outcomes

^a读写寄存器，也就是读写变量，虽然最初与计算机系统中的“寄存器”概念相关，但已慢慢解耦。

Research II (≥ 2017): Replicated Data Types^b



(a) Google Docs



(b) Apache Wave



(c) Wikipedia



(d) L^AT_EX Editor

TODO: +research outcomes

^b复制数据类型，是经典数据类型的分布式版本，如列表，集合，队列等。

Research III (≥ 2020): Distributed Transactions^c
TODO: +research outcomes TODO: +logos

^c分布式事务。每个事务由一组操作构成，这些操作要么全成功，要么全不成功。

Overview of the Work on UNISTORE

UNISTORE: A fault-tolerant marriage of causal and strong consistency

Manuel Bravo

Alexey Gotsman

Borja de Régil

Hengfeng Wei *

IMDEA Software Institute

Nanjing University

ATC'2021 (CCF A)

UNISTORE is the first **fault-tolerant** and scalable **transactional** data store that **combines causal and strong consistency**.

Overview of the Work on UNISTORE

Partial Order-Restrictions Consistency (PoR consistency)

$CC < PoR < SER$

CC: CAUSALCONSISTENCY; SER: SERIALIZABILITY

Key Challenges (I): Ensure liveness in presence of faults

Key Challenges (II): Provide rigorous correctness proof

Overview of the Work on UNISTORE

UNISTORE: A fault-tolerant marriage of causal and strong consistency

Manuel Bravo

Alexey Gotsman

Borja de Régil

Hengfeng Wei *

IMDEA Software Institute

Nanjing University

ATC'2021 (CCF A)

I am fully responsible for the rigorous correctness proof:

- ▶ Finished a proof of 20 pages contained in the arXiv version
- ▶ Identified several nontrivial bugs in the early versions of the protocol^d

^dOne of these bugs also exists in the well-known Granola protocol proposed by James Cowling and Barbara Liskov, something that had gone unnoticed for 10 years.

What is UNISTORE?

UNISTORE is a **fast**, **scalable**, and **fault-tolerant**
transactional distributed key-value store
that supports a **combination of weak and strong consistency**.

What is UNISTORE?

UNISTORE is a **fast**, **scalable**, and **fault-tolerant**
transactional distributed key-value store
that supports a **combination of weak and strong consistency**.


Weak consistency: CAUSALCONSISTENCY

Strong consistency: SERIALIZABILITY

CAUSAL CONSISTENCY and SERIALIZABILITY

Why UNI-?

Weak consistency: low latency, high availability



Strong consistency: easy to preserve critical application invariants

Why UNI-?

DEPOSIT

WITHDRAW

QUERY

INTEREST



Invariant: $\text{balance} \geq 0$

Why UNI-?

DEPOSIT

WITHDRAW

QUERY

INTEREST



Invariant: $\text{balance} \geq 0$

Causal consistency allows two concurrent WITHDRAW to execute without knowing each other.

Why UNI-?

DEPOSIT

WITHDRAW

QUERY

INTEREST



Invariant: $\text{balance} \geq 0$

Causal consistency allows two concurrent WITHDRAW to execute without knowing each other.

Only WITHDRAW needs to use strong consistency.

Consistency Model of UNISTORE

UNISTORE implements a transactional variant of
Partial Order-Restrictions (PoR) consistency [Li@ACT'2018]

- (I) transactional causal consistency by default
- (II) to specify conflicting transactions under strong consistency

Consistency Model of UniSTORE

Definition (Session Order)

A transaction t_1 precedes a transaction t_2 in the **session order**, denoted $t_1 \xrightarrow{so} t_2$, if they are executed by the same client and t_1 is executed before t_2 .

Definition (Conflict Relation)

The **conflict relation**, denoted \bowtie , between transactions is a symmetric relation.

$$t_1 \bowtie t_2 \iff t_2 \bowtie t_1.$$

Consistency Model of UniSTORE

Definition (PoR)

A set of transactions $T \triangleq T_{causal} \uplus T_{strong}$ committed by UniSTORE satisfies **PoR** if there exists a **causal order** \prec on T such that

Consistency Model of UNiSTORE

Definition (PoR)

A set of transactions $T \triangleq T_{causal} \uplus T_{strong}$ committed by UNiSTORE satisfies **PoR** if there exists a **causal order** \prec on T such that

CAUSALITY: ‘ \prec ’ is a partial order and so $\subseteq \prec$.

Consistency Model of UNISTORE

Definition (PoR)

A set of transactions $T \triangleq T_{causal} \uplus T_{strong}$ committed by UNISTORE satisfies **PoR** if there exists a **causal order** \prec on T such that

CAUSALITY: ‘ \prec ’ is a partial order and so $\subseteq \prec$.

CONFLICTORDERING: $\forall t_1, t_2 \in T_{strong}. t_1 \bowtie t_2 \implies t_1 \prec t_2 \vee t_2 \prec t_1$.

Consistency Model of UNISTORE

Definition (PoR)

A set of transactions $T \triangleq T_{causal} \uplus T_{strong}$ committed by UNISTORE satisfies **PoR** if there exists a **causal order** \prec on T such that

CAUSALITY: ‘ \prec ’ is a partial order and $so \subseteq \prec$.

CONFLICTORDERING: $\forall t_1, t_2 \in T_{strong}. t_1 \bowtie t_2 \implies t_1 \prec t_2 \vee t_2 \prec t_1$.

EVENTUALVISIBILITY: A transaction $t \in T$ that is either **strong** or **originates at a correct data center** eventually become **visible** at all **correct** data centers: from some point on, t precedes **in** \prec all transactions issued at correct data centers.

Consistency Model of UNISTORE

Definition (PoR)

A set of transactions $T \triangleq T_{causal} \uplus T_{strong}$ committed by UNISTORE satisfies **PoR** if there exists a **causal order** \prec on T such that

CAUSALITY: ‘ \prec ’ is a partial order and so $\subseteq \prec$.

CONFLICTORDERING: $\forall t_1, t_2 \in T_{strong}. t_1 \bowtie t_2 \implies t_1 \prec t_2 \vee t_2 \prec t_1$.

EVENTUALVISIBILITY: A transaction $t \in T$ that is either **strong** or **originates at a correct data center** eventually become **visible** at all **correct** data centers: from some point on, t precedes **in** \prec all transactions issued at correct data centers.

RETVAL: $\text{INTRETVAL} \wedge \text{EXTRETVAL}$

Consistency Model of UNISTORE

Consider a read r from key k in a transaction t .

INTRETVAL : read from the latest update on k preceding r in t

$$\text{RETV} = \text{INTRETVAL} \wedge \text{EXTRETVAL}$$

EXTRETVAL : read from the last update on k
of the latest transaction (in an order consistent with \prec) preceding t

Consistency Model of UNiSTORE

DEPOSIT **WITHDRAW** QUERY INTEREST



Invariant: $\text{balance} \geq 0$

Declaring that strong transactions
including **WITHDRAW** on the same account conflict.

Design Challenges of UNISTORE


How to satisfy **liveness** (**EVENTUAL VISIBILITY**) despite failures?



A transaction $t \in T$ that is either **strong** or **originates at a correct data center** eventually become **visible** at all **correct** data centers.


Design Challenge of UNISTORE (I)

Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Design Challenge of UNISTORE (I)


Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Transaction t_2 (at correct data center d_2)
may never become visible at correct data center d_3 .


Fault-tolerance of UNISTORE (I)

Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Fault-tolerance of UNISTORE (I)


Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Data center d_2 need to **forward** causal transactions
to other data centers.


Design Challenge of UNISTORE (II)

Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Design Challenge of UNISTORE (II)


Data center d_1 crashes
before t_1 is replicated to correct data center d_3 .



Transaction t_2 will never be visible at d_3 .
No transaction t_3 conflicting with t_2 can commit
(by CONFLICTORDERING).


Fault-tolerance of UNISTORE (II)

UNISTORE ensures that before a strong transaction commits, all its causal dependencies are **uniform**, i.e., will eventually become visible at all correct data centers.



Fault-tolerance of UNISTORE (II)

UNISTORE ensures that before a strong transaction commits, all its causal dependencies are **uniform**, i.e., will eventually become visible at all correct data centers.




Transaction t_1 will eventually be visible at d_3 .

Transaction t_2 will eventually be visible at d_3 .

Transaction t_3 may be committed at d_3 .

Performance of UNISTORE


Causal transactions remain highly-available, i.e., committed locally.



A strong transaction may have to **wait** for some of its dependencies to become uniform before committing.

Performance of UNISTORE

Causal transactions remain highly-available, i.e., committed locally.




A strong transaction may have to **wait** for some of its dependencies to become uniform before committing.

However, this may cost too much.

Performance of UNISTORE

UNISTORE makes a remote causal transaction visible to clients
only **after** it is **uniform**.


Causal transactions are executed on an (almost)
uniform snapshot that may be slightly in the past.



Performance of UNISTORE

UNISTORE makes a remote causal transaction visible to clients
only **after** it is **uniform**.

Causal transactions are executed on an (almost)
uniform snapshot that may be slightly in the past.



A strong transaction only needs to wait for causal transactions
originating at the **local** data center to become uniform.


Scalability of UNISTORE

UNISTORE scales horizontally,
i.e., with the number of machines (partitions) in each data center.

System Model

$\mathcal{D} = \{1, \dots, D\}$: the set of data centers


$\mathcal{P} = \{1, \dots, N\}$: the set of (logical) partitions



p_d^m : the **replica** of partition m at data center d

System Model

$D = 2f + 1$ and $\leq f$ data centers may fail



Any two replicas are connected by a reliable FIFO channel.
Messages between correct data centers will eventually be delivered.

System Model

Replicas have loosely synchronized physical clocks.



The correctness of UNiSTORE does *not* depend on the precision of clock synchronization.

Fault-tolerant Causal Consistency Protocol

Requirement: Tracking Uniformity

UNISTORE makes a remote causal transaction visible to clients
only **after it is uniform**.

Requirement: Tracking Uniformity

UNISTORE makes a remote causal transaction visible to clients only **after it is uniform**.

Definition (Uniform)

A transaction is **uniform** if both the transaction and its causal dependencies are guaranteed to be eventually replicated at all correct data centers.

Requirement: Tracking Uniformity

UNISTORE makes a remote causal transaction visible to clients only **after it is uniform**.

Definition (Uniform)

A transaction is **uniform** if both the transaction and its causal dependencies are guaranteed to be eventually replicated at all correct data centers.

A transaction is considered **uniform** once it is visible at $f + 1$ data centers.

Metadata for Causal Transactions

Each transaction is tagged with a commit vector *commitVec*.

$$\textit{commitVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

For a transaction originating at data center d ,
we call $\textit{commitVec}[d]$ its **local timestamp**.

Metadata for Causal Transactions

Each transaction is tagged with a commit vector *commitVec*.

$$\textit{commitVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

For a transaction originating at data center d ,
we call $\textit{commitVec}[d]$ its **local timestamp**.

Commit vectors are sent to sibling replicas
via **replication and forwarding**.

Metadata for Causal Transactions

Each replica p_d^m maintains the following three vectors:

$$\text{knownVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

$$\text{stableVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Metadata for Causal Transactions

$$\text{knownVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Property (Property of knownVec)

For each data center i ,
the **replica** p_d^m stores the updates to partition m
by transactions originating at i with local timestamps $\leq \text{knownVec}[i]$.

Metadata for Causal Transactions

$$\text{stableVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Property (Property of stableVec)

For each data center i ,
the **data center d** stores the updates
by transactions originating at i with local timestamps $\leq \text{stableVec}[i]$.

Metadata for Causal Transactions

```
1: function BROADCAST_VECS()  
2:   send KNOWNVEC_LOCAL( $m$ , knownVec) to  $p_d^l, l \in \mathcal{P}$   
3:   send STABLEVEC( $d$ , stableVec) to  $p_i^m, i \in \mathcal{D}$   
4:   send KNOWNVEC_GLOBAL( $d$ , knownVec) to  $p_i^m, i \in \mathcal{D}$ 
```

$$\text{stableVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

```
5: when received KNOWNVEC_LOCAL( $l$ , knownVec)  
6:   localMatrix[ $l$ ]  $\leftarrow$  knownVec  
7:   for  $i \in \mathcal{D}$  do  
8:     stableVec[ $i$ ]  $\leftarrow$   $\min\{\text{localMatrix}[n][i] \mid n \in \mathcal{P}\}$   
9:     stableVec[strong]  $\leftarrow$   $\min\{\text{localMatrix}[n][\text{strong}] \mid n \in \mathcal{P}\}$ 
```


Metadata for Causal Transactions

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Property (Property of uniformVec)

All update transactions originating at i
with local timestamps $\leq \text{uniformVec}[i]$
are replicated at $f+1$ data centers including d .

Metadata for Causal Transactions

```
1: function BROADCAST_VECS()  
2:   send KNOWNVEC_LOCAL( $m$ , knownVec) to  $p_d^l, l \in \mathcal{P}$   
3:   send STABLEVEC( $d$ , stableVec) to  $p_i^m, i \in \mathcal{D}$   
4:   send KNOWNVEC_GLOBAL( $d$ , knownVec) to  $p_i^m, i \in \mathcal{D}$ 
```

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

```
10: when received STABLEVEC( $i$ , stableVec)  
11:   stableMatrix[ $i$ ]  $\leftarrow$  stableVec  
12:    $G \leftarrow$  all groups with  $f+1$  replicas that include  $p_d^m$   
13:   for  $j \in \mathcal{D}$  do  
14:     var  $ts \leftarrow \max\{\min\{\text{stableMatrix}[h][j] \mid h \in g\} \mid g \in G\}$   
15:     uniformVec[ $j$ ]  $\leftarrow \max\{\text{uniformVec}[j], ts\}$ 
```

Metadata for Causal Transactions

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Lemma

*All update transactions
with **commit vectors** \leq **uniformVec** are uniform.*

Metadata for Causal Transactions

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

Lemma

*All update transactions
with **commit vectors** \leq **uniformVec** are uniform.*

UNISTORE makes a remote causal transaction visible to clients
only **after it is uniform.**

Causal Consistency Protocol: Start

pastVec : causal past of client

- 1: **function** START()
- 2: $p \leftarrow$ a random partition in data center d
- 3: $\langle \text{tid}, \text{snapVec} \rangle \leftarrow$ **send** START_TX(pastVec) **to** p
- 4: $\text{pastVec} \leftarrow \text{snapVec}$
- 5: **return** tid

$\forall i \in \mathcal{D} \setminus \{d\}$, all transactions originating at i
with local timestamps $\leq \text{pastVec}[i]$ are already uniform.

Causal Consistency Protocol: Start

Causal transactions are executed on an (almost) **uniform snapshot**.

```
1: function START_TX( $V$ )
2:   for  $i \in \mathcal{D} \setminus \{d\}$  do
3:      $\text{uniformVec}[i] \leftarrow \max\{V[i], \text{uniformVec}[i]\}$ 
4:   var  $tid \leftarrow \text{generate\_tid}()$ 
5:    $\text{snapVec}[tid] \leftarrow \text{uniformVec}$ 
6:    $\text{snapVec}[tid][d] \leftarrow \max\{V[d], \text{uniformVec}[d]\}$ 
7:    $\text{snapVec}[tid][\text{strong}] \leftarrow \max\{V[\text{strong}], \text{stableVec}[\text{strong}]\}$ 
8:   return  $\langle tid, \text{snapVec}[tid] \rangle$ 
```

$\text{snapVec}[tid][d]$ ensures “read-your-writes”.

Causal Consistency Protocol: Update

```
11: function UPDATE( $k, v$ )  
12:    $\_ \leftarrow$  send DO_UPDATE( $tid, k, v$ ) to  $p$   
13:   return ok
```

```
17: function DO_UPDATE( $tid, k, v$ )  
18:   var  $l \leftarrow$  partition( $k$ )  
19:   wbuff[ $tid$ ][ $l$ ][ $k$ ]  $\leftarrow v$   
20:   rset[ $tid$ ][ $l$ ]  $\leftarrow$  rset[ $tid$ ][ $l$ ]  $\cup \{k\}$   
21:   return ok
```

wbuff[tid][l] : buffer for the latest local update on each key

Causal Consistency Protocol: Read

```
6: function READ( $k$ )
7:    $\langle v, c \rangle \leftarrow \text{send DO\_READ}(\text{tid}, k, \text{lc})$  to  $p$ 
8:   if  $c \neq \perp$  then
9:      $\text{lc} \leftarrow \max\{\text{lc}, c\}$ 
10:  return  $v$ 
```

```
9: function DO_READ( $\text{tid}, k, c$ )
10:   $\text{lc} \leftarrow \max\{\text{lc}, c\}$ 
11:  var  $l \leftarrow \text{partition}(k)$ 
12:  if  $\text{wbuff}[\text{tid}][l][k] \neq \perp$  then
13:    return  $\langle \text{wbuff}[\text{tid}][l][k], \perp \rangle$ 
14:   $\langle v, c \rangle \leftarrow \text{send READ\_KEY}(\text{snapVec}[\text{tid}], k)$  to  $p_d^l$ 
15:   $\text{rset}[\text{tid}][l] \leftarrow \text{rset}[\text{tid}][l] \cup \{k\}$ 
16:  return  $\langle v, c \rangle$ 
```


Causal Consistency Protocol: Read

Causal transactions are executed on an (almost) **uniform snapshot**.

- 1: **when received** `READ_KEY(snapVec, k)` **from** *p*
- 2: **for** $i \in \mathcal{D} \setminus \{d\}$ **do**
- 3: `uniformVec[i] \leftarrow max{snapVec[i], uniformVec[i]}`
- 4: **wait until** `knownVec[d] \geq snapVec[d] \wedge knownVec[strong] \geq snapVec[strong]`
- 5: $\langle v, \text{commitVec}, c \rangle \leftarrow \text{snapshot}(\text{opLog}[k], \text{snapVec})$ \triangleright returns the latest *commitVec* (in terms of Lamport clock order in Definition 50) such that *commitVec* \leq *snapVec*
- 6: **send** $\langle v, c \rangle$ **to** *p*

wait : ensure that it is as up-to-date as required by the snapshot

Causal Consistency Protocol: Commit

```
14: function COMMIT_CAUSAL_TX()  
15:    $\langle vc, c \rangle \leftarrow \text{send COMMIT\_CAUSAL}(\text{tid}, lc) \text{ to } p$   
16:    $\text{pastVec} \leftarrow vc$   
17:    $lc \leftarrow c$   
18:   return ok
```

Causal Consistency Protocol: Commit

Read-only transactions returns immediately.

```
22: function COMMIT_CAUSAL( $tid, c$ )
23:    $lc \leftarrow \max\{lc, c\} + 1$ 
24:   if  $\forall l \in \mathcal{P}. wbuff[tid][l] = \emptyset$  then
25:     return  $\langle snapVec[tid], lc \rangle$ 
26:   var  $commitVec \leftarrow snapVec[tid]$ 
27:   send PREPARE( $tid, wbuff[tid][l], snapVec[tid]$ ) to  $p_d^l, l \in \mathcal{P}$ 
28:   for all  $l \in \mathcal{P}$  do
29:     wait receive PREPARE_ACK( $tid, ts$ ) from  $p_d^l$ 
30:      $commitVec[d] \leftarrow \max\{commitVec[d], ts\}$ 
31:   send COMMIT( $tid, commitVec, lc$ ) to  $p_d^l, l \in \mathcal{P}$ 
32:   return  $\langle commitVec, lc \rangle$ 
```

2PC protocol for update transactions

Causal Consistency Protocol: Commit

ts : **prepare timestamp** from its local clock

```
7: when received PREPARE( $tid, wbuff, snapVec$ ) from  $p$ 
8:   for  $i \in \mathcal{D} \setminus \{d\}$  do
9:      $uniformVec[i] \leftarrow \max\{snapVec[i], uniformVec[i]\}$ 
10:  var  $ts \leftarrow \text{clock}$ 
11:   $preparedCausal \leftarrow preparedCausal \cup \{\langle tid, wbuff, ts \rangle\}$ 
12:  send PREPARE_ACK( $tid, ts$ ) to  $p$ 
```

Causal Consistency Protocol: Commit

wait : ensure that its local clock is up-to-date

```
13: when received COMMIT(tid, commitVec, c)
14:   wait until clock  $\geq$  commitVec[d]
15:    $\langle tid, wbuff, - \rangle \leftarrow \text{find}(tid, \text{preparedCausal})$ 
16:   preparedCausal  $\leftarrow$  preparedCausal  $\setminus \{ \langle tid, -, - \rangle \}$ 
17:   for all  $\langle k, v \rangle \in wbuff$  do
18:      $\text{opLog}[k] \leftarrow \text{opLog}[k] \cdot \langle v, \text{commitVec}, c \rangle$ 
19:   committedCausal[d]  $\leftarrow$  committedCausal[d]  $\cup \{ \langle tid, wbuff, \text{commitVec}, c \rangle \}$ 
```

committedCausal[*d*] : for replication

Causal Consistency Protocol: Replication

Property (Property of knownVec)

For each data center i ,
the replica p_d^m stores the updates to partition m
by transactions originating at i with local timestamps $\leq \text{knownVec}[i]$.

```
1: function PROPAGATE_LOCAL_TXS()  
2:   if preparedCausal =  $\emptyset$  then  
3:     knownVec[d]  $\leftarrow$  clock  
4:   else  
5:     knownVec[d]  $\leftarrow$  min{ts |  $\langle -, -, ts \rangle \in \text{preparedCausal}$ } - 1  
6:   var txs  $\leftarrow$  { $\langle -, -, \text{commitVec}, c \rangle \in \text{committedCausal}[d] \mid \text{commitVec}[d] \leq \text{knownVec}[d]$ }  
7:   if txs  $\neq \emptyset$  then  
8:     send REPLICATE( $d, txs$ ) to  $p_i^m, i \in \mathcal{D} \setminus \{d\}$   
9:     committedCausal[d]  $\leftarrow$  committedCausal[d]  $\setminus$  txs  
10:  else  
11:    send HEARTBEAT( $d, \text{knownVec}[d]$ ) to  $p_i^m, i \in \mathcal{D} \setminus \{d\}$ 
```

HEARTBEAT : for liveness

Adding Strong Transactions

Requirement: CONFLICTORDERING

$$\forall t_1, t_2 \in T_{strong}. t_1 \bowtie t_2 \implies t_1 \prec t_2 \vee t_2 \prec t_1.$$

Each strong transaction is assigned a scalar **strong timestamp**.

$$commitVec \in [\mathcal{D} \cup \{strong\} \rightarrow \mathbb{N}]$$

Metadata for Strong Transactions

$$\text{knownVec} \in [\mathcal{D} \cup \{\text{strong}\}] \rightarrow \mathbb{N}$$

Property (Property of $\text{knownVec}[\text{strong}]$)

Replica p_d^m stores the updates to m by all strong transactions with $\text{commitVec}[\text{strong}] \leq \text{knownVec}[\text{strong}]$.

Metadata for Strong Transactions

$$\text{stableVec} \in [\mathcal{D} \cup \{\text{strong}\}] \rightarrow \mathbb{N}$$

```
5: when received KNOWNVEC_LOCAL( $l, \text{knownVec}$ )  
6:    $\text{localMatrix}[l] \leftarrow \text{knownVec}$   
7:   for  $i \in \mathcal{D}$  do  
8:      $\text{stableVec}[i] \leftarrow \min\{\text{localMatrix}[n][i] \mid n \in \mathcal{P}\}$   
9:    $\text{stableVec}[\text{strong}] \leftarrow \min\{\text{localMatrix}[n][\text{strong}] \mid n \in \mathcal{P}\}$ 
```

Property (Property of $\text{stableVec}[\text{strong}]$)

Data center d stores the updates by all strong transactions
with $\text{commitVec}[\text{strong}] \leq \text{knownVec}[\text{strong}]$.

Metadata for Strong Transactions

$$\text{uniformVec} \in [\mathcal{D} \rightarrow \mathbb{N}]$$

```
10: when received STABLEVEC( $i, \text{stableVec}$ )  
11:    $\text{stableMatrix}[i] \leftarrow \text{stableVec}$   
12:    $G \leftarrow$  all groups with  $f + 1$  replicas that include  $p_d^m$   
13:   for  $j \in \mathcal{D}$  do  
14:     var  $ts \leftarrow \max\{\min\{\text{stableMatrix}[h][j] \mid h \in g\} \mid g \in G\}$   
15:      $\text{uniformVec}[j] \leftarrow \max\{\text{uniformVec}[j], ts\}$ 
```

The commit protocol for strong transactions
guarantees their uniformity.

Strong Consistency Protocol: Commit

```
1: function COMMIT_STRONG( $tid, c$ )
2:   UNIFORM_BARRIER( $\text{snapVec}[tid]$ )
3:    $\langle d, vc, c \rangle \leftarrow \text{CERTIFY}(tid, \text{wbuff}[tid], \text{rset}[tid], \text{snapVec}[tid], c)$ 
4:    $lc \leftarrow \max\{lc, c\} + 1$ 
5:   return  $\langle d, vc, lc \rangle$ 
```

A strong transaction only needs to wait for causal transactions originating at the **local** data center to become uniform.

```
20: function UNIFORM_BARRIER( $V, c$ )
21:    $lc \leftarrow \max\{lc, c\} + 1$ 
22:   wait until  $\text{uniformVec}[d] \geq V[d]$ 
23:   return  $lc$ 
```

Strong Consistency Protocol: Commit

```
1: function COMMIT_STRONG(tid, c)  
2:   UNIFORM_BARRIER(snapVec[tid])  
3:    $\langle d, vc, c \rangle \leftarrow \text{CERTIFY}(tid, wbuff[tid], rset[tid], snapVec[tid], c)$   
4:    $lc \leftarrow \max\{lc, c\} + 1$   
5:   return  $\langle d, vc, lc \rangle$ 
```

$$\langle d \in \{\text{COMMIT}, \text{ABORT}\}, vc \rangle \leftarrow \text{CERTIFY}(t)$$

Strong Consistency Protocol: Commit

```
1: function COMMIT_STRONG(tid, c)  
2:   UNIFORM_BARRIER(snapVec[tid])  
3:    $\langle d, vc, c \rangle \leftarrow$  CERTIFY(tid, wbuff[tid], rset[tid], snapVec[tid], c)  
4:    $lc \leftarrow \max\{lc, c\} + 1$   
5:   return  $\langle d, vc, lc \rangle$ 
```

$$\langle d \in \{\text{COMMIT}, \text{ABORT}\}, vc \rangle \leftarrow \text{CERTIFY}(t)$$

Multi-Shot Distributed Transaction Commit

Gregory Chockler

Royal Holloway, University of London, UK

Alexey Gotsman¹

IMDEA Software Institute, Madrid, Spain

White-Box Atomic Multicast

Alexey Gotsman
IMDEA Software Institute

Anatole Lefort
Télécom SudParis

Gregory Chockler
Royal Holloway, University of London

2PC across partitions + Paxos among replicas of each partition
uses white-box optimizations that minimize the commit latency


Strong Consistency Protocol: Deliver

```
6: upon DELIVER_UPDATES( $W$ )
7:   for  $\langle k, v, \text{commitVec}, c \rangle \in W$  in  $\text{commitVec}[\text{strong}]$  order do
8:      $\text{opLog}[k] \leftarrow \text{opLog}[k] \cdot \langle v, \text{commitVec}, c \rangle$ 
9:      $\text{knownVec}[\text{strong}] \leftarrow \text{commitVec}[\text{strong}]$ 
```

Evaluation

Performance of UNISTORE


Throughput: 5% and 259% higher than REDBLUE and STRONG



RUBiS benchmark: throughput vs. average latency.

Latency: 24ms vs. 32ms of REDBLUE and 162ms of STRONG

Scalability of UNISTORE



Scalability when varying the ratio of strong transactions.

UNISTORE is able to scales almost linearly.

Evaluation

For more evaluations, please refer to the paper.

Conclusion

UNISTORE is a **fast**, **scalable**, and **fault-tolerant**
transactional distributed key-value store
that supports a **combination of weak and strong consistency**.

Conclusion

UNISTORE is a **fast**, **scalable**, and **fault-tolerant**
transactional distributed key-value store
that supports a **combination of weak and strong consistency**.

“We expect the key ideas in UNISTORE to pave the way
for practical systems that combine causal and strong consistency.”

总结

魏恒峰 (hfwei@nju.edu.cn)

聘期合同要求	工作情况
教学: 承担一门课程	问题求解课程 五个学期; 共 164 学时 (2019 级本科生“我心目中的好课程”)
科研: 4-6 篇高水平论文	发表 3 篇 (含 1 篇短文) 在审 4 篇 (2017 年 CCF 优秀博士学位论文奖)
人才培养	负责或协助指导学生 9 人次 (学术积累: 组织 TLA ⁺ 与 Coq 讨论班)
主持/参与 多个基金项目	主持 1 项; 参与 1 项 个人可支配总经费 75 万元



Hengfeng Wei (hfwei@nju.edu.cn)

